CSE 590 **Introduction to Machine Learning** Spring 2022

Homework No. 3
Due Feb. 25 (11:59pm), 2022

Objectives

1. *Apply various classification algorithms to the movie reviews dataset*
2. *Use k-fold cross validation to identify the parameters that optimize performance (generalization) for each method*
3. *Compare the accuracy and explainability of each method*

Problem #1

For this homework, you will apply the following classification methods to the *movie reviews classification data* (available in Blackboard)

1. Multinomial Naïve Bayes
2. Random Forest
3. Gradient Boosted Regression Trees

- Apply 4-fold cross-validation to the provided training data subset to train your classifiers and identify their *optimal parameters*.
- After fixing the classifiers' parameters, apply each method to the provided testing data subset to predict and analyze your results. *Compare the accuracy* obtained during training (average of the cross-validation folds) to those of the test data and comment on the results (overfitting, underfitting, etc.)
- Analyze the results of each method by *inspecting the feature importance* (if applicable) and few misclassified samples.
- Select the best algorithm and justify your choice based on *accuracy*, *explainability*, *time required to train/test*, etc.

What to submit?

- A report that
  - **Describes** your experiments,
  - **Summarizes**, **explains** (using concepts covered in lectures) and **compares** the results (using plots, tables, figures)
  - Identifies the best method for each dataset.
- Do not submit your source code
- Do not submit raw output generated by your code!
- Your report needs to be a single file (MS Word or PDF)
- Your report cannot exceed 10 pages using a font of 12
- Assign numbers to all your figures/tables/plots and use these numbers to reference them in your discussion